

Deterministic SIR Model with Stochastic Cases, with Normalising Flow

```
library(torch)
library(torchflow)
```

Stochastic SIR model

Forward simulation

The following function simulates the stochastic SIR model. The model is defined by the following parameters:

- N is the total population size.
- D is the duration of the infection.
- R_0 is the basic reproduction number.
- $\gamma = 1/D$ is the rate of recovery.
- $\beta = R_0\gamma$ is the rate of infection.
- ϵ is the rate of new infections from outside the population.
- α is the rate of reporting cases.
- S_1, I_1, R_1 are the initial number of susceptible, infected, and recovered individuals

At time t , the number of new infections is given by the binomial distribution:

$$I_t \sim \text{Binomial}(S_{t-1}, \lambda_t)$$

where $\lambda_t = 1 - \exp(-\beta I_{t-1}/N - \epsilon)$. The parameter ϵ is the rate of new infections from outside the population, which ensures that an epidemic can reappear even after it is extinct.

The number of new recoveries is given by the binomial distribution:

$$R_t \sim \text{Binomial}(I_{t-1}, \gamma)$$

where γ is the rate of recovery.

The number of new cases is given by the binomial distribution:

$$C_t \sim \text{Binomial}(I_t, \alpha)$$

```
forward_simulate <- function(
  T,
  N,
  alpha,
  beta,
  gamma,
  epsilon,
  S_1,
  I_1,
  R_1
) {
  n <- length(alpha)
```

```

S <- matrix(0, n, T)
I <- matrix(0, n, T)
R <- matrix(0, n, T)
cases <- matrix(0, n, T)

S[, 1] <- round(N * S_1)
I[, 1] <- round(N * I_1)
R[, 1] <- round(N * R_1)

for (t in 2 : T) {
  lambda <- 1 - exp(-beta * I[, t - 1] / N - epsilon)
  new_infections <- rbinom(n, S[, t - 1], lambda)
  new_recoveries <- rbinom(n, I[, t - 1], gamma)
  new_cases <- rbinom(n, new_infections, alpha)

  S[, t] <- S[, t - 1] - new_infections
  I[, t] <- I[, t - 1] + new_infections - new_recoveries
  R[, t] <- R[, t - 1] + new_recoveries
  cases[, t] <- new_cases
}

list(
  S = S / N,
  I = I / N,
  R = R / N,
  cases = cases / N
)
}

```

We normalise S, I, R, and cases to be between 0 and 1, a proportion of the total population.

Here is an example of a forward simulation:


```

set.seed(20250720)
example_simulation <- forward_simulate(
  T = 100,
  N = 300,
  alpha = 0.25,
  beta = 0.5,
  gamma = 0.2,
  epsilon = 0.0001,
  S_1 = 1,
  I_1 = 0,
  R_1 = 0
)

plot(1 : 100, example_simulation$S[1, ], type = "l", xlab = 'Time', col = 'blue', ylim = c(0, 1))
lines(1 : 100, example_simulation$I[1, ], col = 'red')
lines(1 : 100, example_simulation$R[1, ], col = 'green')
lines(1 : 100, example_simulation$cases[1, ], type = "h", col = 'black')
legend(
  'topright', c('S', 'I', 'R', 'Cases'),
  col = c('blue', 'red', 'green', 'black'),
  lty = 1, bg = 'white'
)

```

)



Here we generate from a prior distribution as follows:

- R_0 is uniformly distributed between 1.01 and 3.
- D is uniformly distributed between 2 and 6.
- ϵ is uniformly distributed between $1e-4$ and $2 * 1e-4$.
- α is uniformly distributed between 0.2 and 0.3.

Once these are generated, we simulate a single trajectory of the model.

```
generate <- function(
  n,
  T,
  N,
  R_0_min = 1.01,
  R_0_max = 3,
  D_min = 2,
  D_max = 6,
  epsilon_min = 1e-4,
  epsilon_max = 2 * 1e-4,
  alpha_min = 0.2,
  alpha_max = 0.3
) {
  R_0 <- runif(n, R_0_min, R_0_max)
  D <- runif(n, D_min, D_max)
  epsilon <- runif(n, epsilon_min, epsilon_max)
```

```

alpha <- runif(n, alpha_min, alpha_max)
gamma <- 1 / D
beta <- R_0 * gamma

simulation <- forward_simulate(
  T,
  N,
  alpha,
  beta,
  gamma,
  epsilon,
  1,
  0,
  0
)

c(simulation, list(
  R_0 = R_0,
  D = D,
  alpha = alpha,
  beta = beta,
  gamma = gamma,
  epsilon = epsilon
))
}

```

We generate four example datasets:

```

T <- 100
N <- 300
example_data <- generate(4, T = T, N = N)
str(example_data)

## List of 10
## $ S      : num [1:4, 1:100] 1 1 1 1 1 1 1 1 1 1 ...
## $ I      : num [1:4, 1:100] 0 0 0 0 0 0 0 0 0 0 ...
## $ R      : num [1:4, 1:100] 0 0 0 0 0 0 0 0 0 0 ...
## $ cases  : num [1:4, 1:100] 0 0 0 0 0 0 0 0 0 0 ...
## $ R_0     : num [1:4] 1.06 1.96 2.47 1.69
## $ D      : num [1:4] 2.4 3.16 3.54 3.66
## $ alpha   : num [1:4] 0.21 0.207 0.251 0.24
## $ beta    : num [1:4] 0.442 0.622 0.697 0.461
## $ gamma   : num [1:4] 0.416 0.316 0.282 0.273
## $ epsilon: num [1:4] 0.000145 0.000177 0.000121 0.000168

print(example_data[c('gamma', 'beta', 'alpha', 'epsilon')])


## $gamma
## [1] 0.4161333 0.3164321 0.2821377 0.2731598
##
## $beta
## [1] 0.4420780 0.6216277 0.6970645 0.4612615
##
## $alpha
## [1] 0.2098382 0.2072343 0.2512811 0.2403872

```

```
##
## $epsilon
## [1] 0.0001453925 0.0001765887 0.0001211146 0.0001678262
```

We can plot the results:

```
withr::with_par(list(mfrow = c(2, 1), mar = c(4, 4, 1, 1)), {
  matplot(t(example_data$I), type = "l", lty = 1, xlab = 'Time', ylab = 'Infections')
  matplot(t(example_data$cases), type = "h", lty = 1, xlab = 'Time', ylab = 'Cases')
})
```



Normalising flow

The following code define a normalising flow model that takes as input the number of cases up to time t . It then approximates the posterior distribution of the $(\alpha, \beta, \gamma, \epsilon, S_t, I_t)$ parameters. More precisely, it approximates a transformation of the parameters:

- $\log \alpha$
- $\log \beta$
- $\log \gamma$
- $\log \epsilon$
- $\text{logit}(S_t)$, where S_t is clamped to be between $1e-5$ and $1 - 1e-5$
- $\text{logit}(I_t)$, where I_t is clamped to be between $1e-5$ and $1 - 1e-5$

The network that summarises the data into a fixed dimensional space is a Gated Recurrent Unit (GRU). This is a type of recurrent neural network that is designed to process sequential data. It is constructed such that, after all GRU layers have been applied, the output at time t is only a function of the data up to time t . This allows us to input varying sequence lengths.

```

device <- if (cuda_is_available()) torch_device("cuda") else torch_device("cpu")

summary_gru <- nn_module(
  "summary_gru",
  initialize = function(n_input_size, n_layers, n_output_size) {
    self$n_layers <- n_layers
    self$gru <- nn_gru(n_input_size, n_output_size, n_layers, batch_first = TRUE)
  },
  forward = function(conditioning) {
    t <- conditioning[, 1, drop = FALSE]
    x <- conditioning[, 2 : conditioning$shape[2], drop = FALSE]

    t_int <- t$to(dtype = torch_int64())
    x_unsqueezed <- x$unsqueeze(-1)

    # h_all contains the hidden state of the GRU at each time point
    h_all <- self$gru(x_unsqueezed)[[1]]

    # h_last is the hidden state at the last time point
    t_int_big <- t_int$unsqueeze(2)$expand(c(-1, -1, h_all$shape[3]))
    h_last <- torch_gather(h_all, 2, t_int_big)
    h_last$squeeze(2)
  }
)

n_summary <- 128L
summary_model <- summary_gru(n_input_size = 1, n_layers = 2, n_output_size = n_summary)
summary_model$to(device = device)

n_params <- 6L
flow_model <- nn_sequential_conditional_flow(
  nn_affine_coupling_block(n_params, n_summary),
  nn_permutation_flow(n_params),
  nn_affine_coupling_block(n_params, n_summary),
  nn_permutation_flow(n_params),
  nn_affine_coupling_block(n_params, n_summary),
  nn_permutation_flow(n_params),
  nn_affine_coupling_block(n_params, n_summary),
  nn_permutation_flow(n_params),
  nn_affine_coupling_block(n_params, n_summary)
)
flow_model$to(device = device)

summarizing_flow_model <- nn_summarizing_conditional_flow(summary_model, flow_model)
summarizing_flow_model$to(device = device)

```

Then we define a function to generate batches of data for training and testing:

```

logit <- function(x) {
  log(x / (1 - x))
}

expit <- function(y) {
  exp(y) / (1 + exp(y))
}

```

```

}

clamp <- function(x, min, max) {
  pmax(pmin(x, max), min)
}

epsilon <- 1e-5

generate_conditional_samples <- function(...) {
  n <- 1024L
  data <- generate(n, T = T, N = N)
  t <- sample.int(T, n, replace = TRUE)

  list(
    target = torch_tensor(cbind(
      log(data$alpha),
      log(data$beta),
      log(data$gamma),
      log(data$epsilon),
      logit(clamp(data$S[cbind(seq_len(n), t)], epsilon, 1 - epsilon)),
      logit(clamp(data$I[cbind(seq_len(n), t)], epsilon, 1 - epsilon))
    ), device = device),
    conditioning = torch_tensor(
      cbind(t, data$cases),
      device = device
    ),
    latent = torch_stack(list(
      data$S,
      data$I,
      data$R
    ), dim = -1)$to(device = device)
  )
}

```

We can calculate the initial test loss:

```

test_set <- generate_conditional_samples()

test_loss <- forward_kl_loss(summarizing_flow_model(test_set$target, test_set$conditioning))
cat('Initial test loss:', test_loss$item(), '\n')

```

Initial test loss: 89.54634

Now we train the model.

```

run_training <- function(n_epochs, lr) {
  system.time(train_conditional_flow(
    summarizing_flow_model,
    generate_conditional_samples,
    n_epochs = n_epochs,
    batch_size = 256L,
    lr = lr,
    after_epoch = function(epoch, ...) {
      if (epoch %% 64L == 0) {
        test_loss <- forward_kl_loss(summarizing_flow_model(test_set$target, test_set$conditioning))

```

```

        cat('Epoch', epoch, '| test loss:', test_loss$item(), '\n')
    }
},
verbose = interactive()
))
}

if (!file.exists('summarizing_flow_model.pt')) {
  run_training(256L, 0.001)
  run_training(1024L, 0.0001)
  run_training(512L, 0.00001)
  torch_save(summarizing_flow_model, 'summarizing_flow_model.pt')
} else {
  summarizing_flow_model <- torch_load('summarizing_flow_model.pt')
  summarizing_flow_model$to(device = device)
}

```

```

## Epoch 64 | test loss: -1.842995
## Epoch 128 | test loss: -1.526629
## Epoch 192 | test loss: -2.377735
## Epoch 256 | test loss: -2.708314
## Epoch 64 | test loss: -4.148572
## Epoch 128 | test loss: -4.985259
## Epoch 192 | test loss: -5.472393
## Epoch 256 | test loss: -5.590963
## Epoch 320 | test loss: -6.118996
## Epoch 384 | test loss: -6.344919
## Epoch 448 | test loss: -6.488613
## Epoch 512 | test loss: -6.64969
## Epoch 576 | test loss: -5.713623
## Epoch 640 | test loss: -6.532421
## Epoch 704 | test loss: -6.661923
## Epoch 768 | test loss: 1.381217
## Epoch 832 | test loss: -6.940806
## Epoch 896 | test loss: -7.0771
## Epoch 960 | test loss: -6.672498
## Epoch 1024 | test loss: -6.124325
## Epoch 64 | test loss: -7.32359
## Epoch 128 | test loss: -7.360783
## Epoch 192 | test loss: -7.387857
## Epoch 256 | test loss: -7.424115
## Epoch 320 | test loss: -7.462863
## Epoch 384 | test loss: -7.474615
## Epoch 448 | test loss: -7.507388
## Epoch 512 | test loss: -7.484043

```

The training is complete.

Results

Let's generate some samples from the approximate posterior.

We will generate samples from the posterior for four entries of the test set.


```

sample_dataset <- function(n_samples, dataset) {
  n_entries <- dataset$conditioning$shape[1]
  samples_tensor <- generate_from_conditional_flow(
    summarizing_flow_model,
    n_samples,
    dataset$conditioning
  )
  samples <- as_array(samples_tensor)
  dimnames(samples) <- list(
    NULL,
    NULL,
    c('alpha', 'beta', 'gamma', 'epsilon', 'S_t', 'I_t')
  )
  samples[, , 'alpha'] <- exp(samples[, , 'alpha'])
  samples[, , 'beta'] <- exp(samples[, , 'beta'])
  samples[, , 'gamma'] <- exp(samples[, , 'gamma'])
  samples[, , 'epsilon'] <- exp(samples[, , 'epsilon'])
  samples[, , 'S_t'] <- expit(samples[, , 'S_t'])
  samples[, , 'I_t'] <- expit(samples[, , 'I_t'])

  forward_samples <- lapply(seq_len(n_entries), function(i) {
    test_t <- as_array(dataset$conditioning[i, 1])[1]
    S_1 <- samples[, i, 'S_t']
    I_1 <- samples[, i, 'I_t']
    R_1 <- 1 - S_1 - I_1
    simulation <- forward_simulate(
      T = T - test_t + 1L,
      N = N,
      alpha = samples[, i, 'alpha'],
      beta = samples[, i, 'beta'],
      gamma = samples[, i, 'gamma'],
      epsilon = samples[, i, 'epsilon'],
      S_1 = S_1,
      I_1 = I_1,
      R_1 = R_1
    )
    simulation$time <- test_t : T
    simulation
  })

  target <- as_array(dataset$target)
  colnames(target) <- c('alpha', 'beta', 'gamma', 'epsilon', 'S', 'I')
  target[, 'alpha'] <- exp(target[, 'alpha'])
  target[, 'beta'] <- exp(target[, 'beta'])
  target[, 'gamma'] <- exp(target[, 'gamma'])
  target[, 'epsilon'] <- exp(target[, 'epsilon'])
  target[, 'S'] <- expit(target[, 'S'])
  target[, 'I'] <- expit(target[, 'I'])

  t <- as_array(dataset$conditioning[, 1])
  cases <- as_array(dataset$conditioning[, 2 : 101])
  latent <- as_array(dataset$latent)

```

```

list(
  samples = samples,
  forward_samples = forward_samples,
  target = target,
  cases = cases,
  latent = latent,
  t = t
)
}

test_t <- as_array(test_set$conditioning[, 1])
example_indices <- c(
  sample(which(test_t < 25), 1),
  sample(which(test_t < 25), 1),
  sample(which(test_t >= 25 & test_t < 50), 1),
  sample(which(test_t >= 25 & test_t < 50), 1),
  sample(which(test_t >= 50 & test_t < 75), 1),
  sample(which(test_t >= 50 & test_t < 75), 1),
  sample(which(test_t >= 75), 1),
  sample(which(test_t >= 75), 1)
)
example_dataset <- list(
  target = test_set$target[example_indices, ],
  conditioning = test_set$conditioning[example_indices, , drop = FALSE],
  latent = test_set$latent[example_indices, , drop = FALSE]
)

example_samples <- sample_dataset(256L, example_dataset)
n_examples <- length(example_indices)

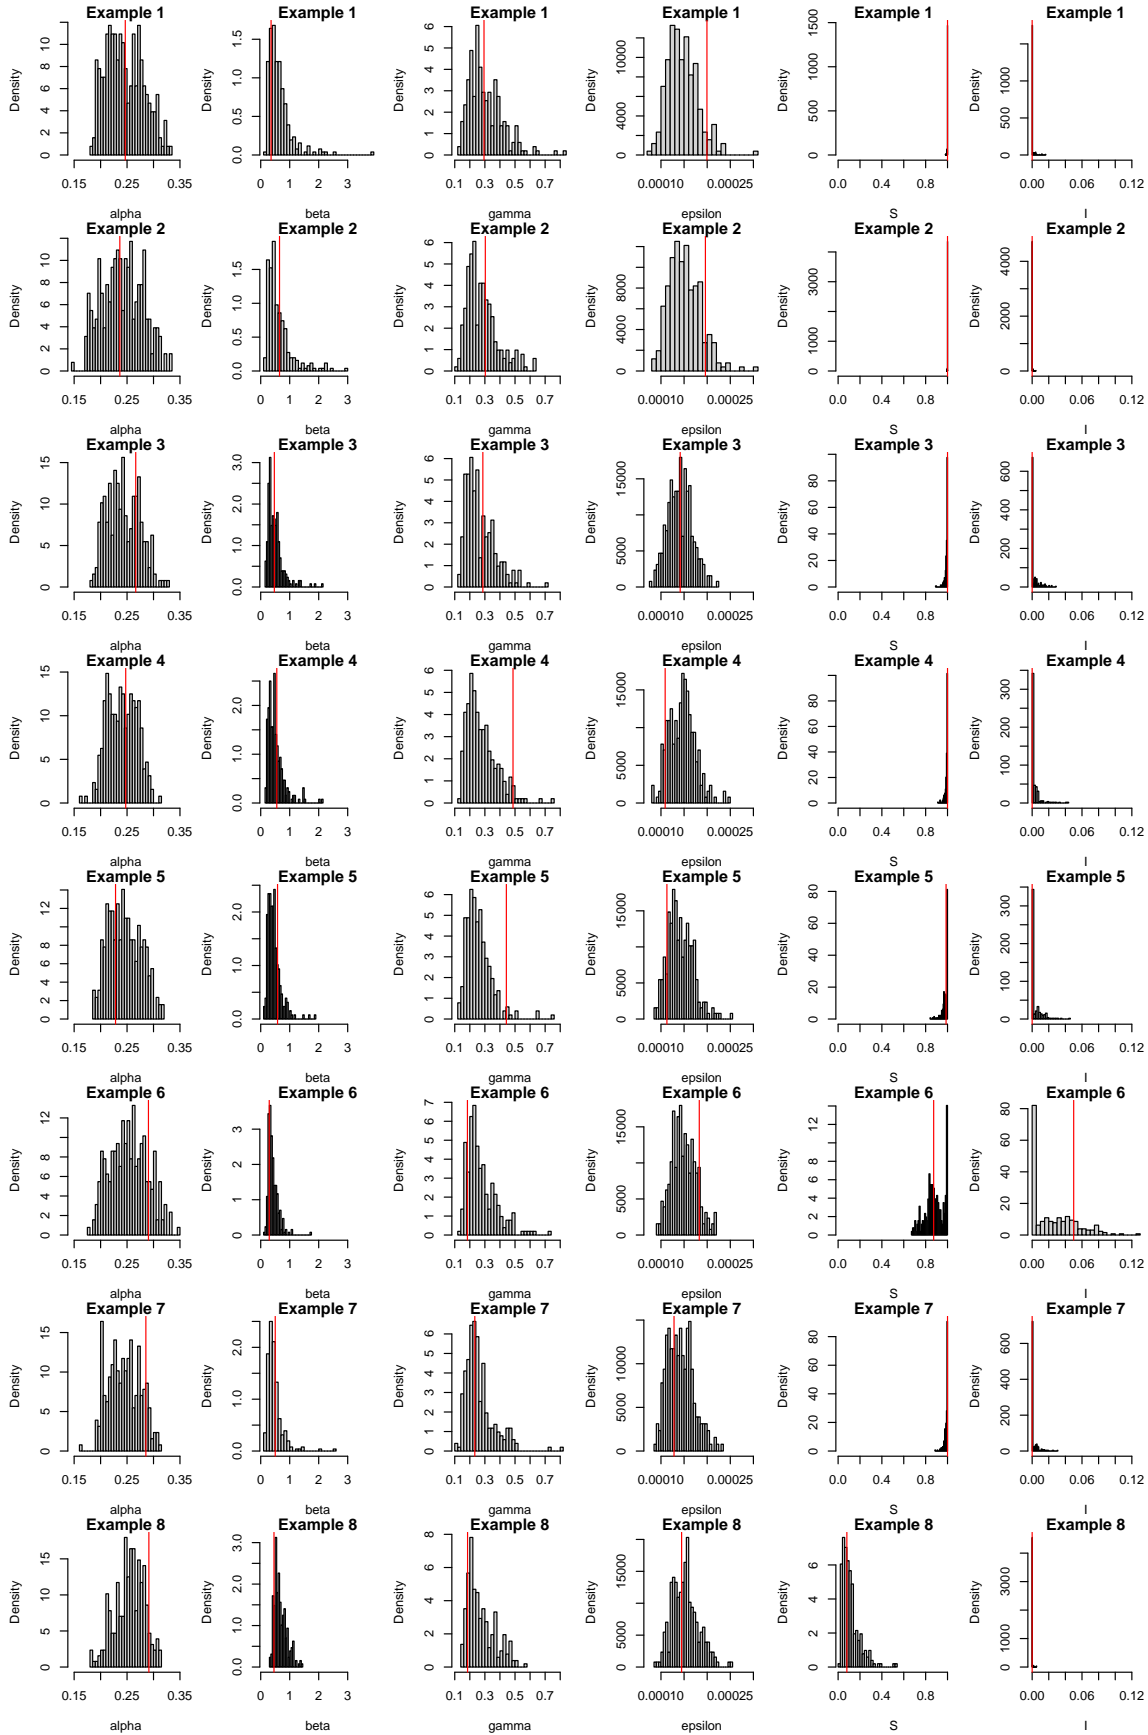
```

Now we will plot the samples and the true values. First, histograms (red lines are the true values):

```

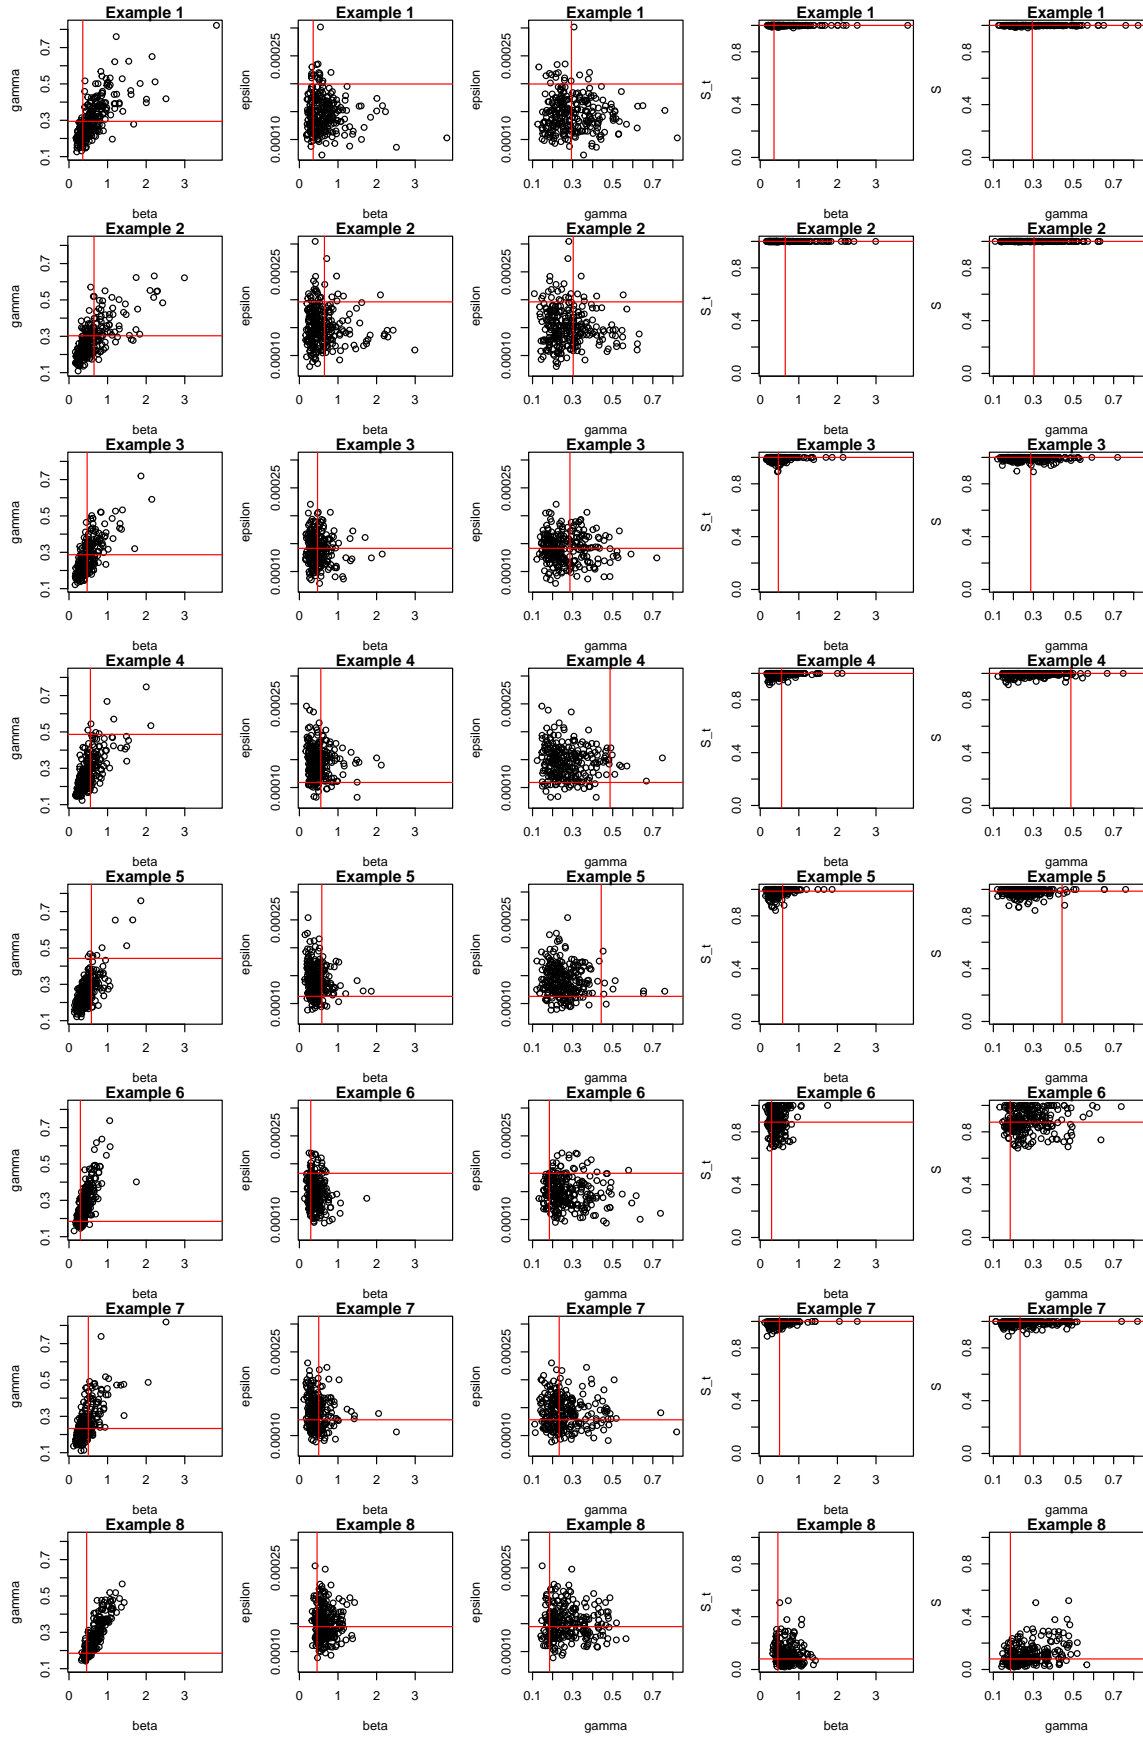
withr::with_par(list(mfrow = c(n_examples, n_params), mar = c(4, 4, 1, 1)), {
  for (i in seq_len(n_examples)) {
    plot_histogram <- function(ind, lab) {
      hist(
        example_samples$samples[, i, ind],
        xlab = lab,
        xlim = range(example_samples$samples[, , ind]),
        main = paste('Example', i),
        freq = FALSE,
        breaks = 30
      )
      abline(v = example_samples$target[i, ind], col = 'red')
    }
    plot_histogram(1, 'alpha')
    plot_histogram(2, 'beta')
    plot_histogram(3, 'gamma')
    plot_histogram(4, 'epsilon')
    plot_histogram(5, 'S')
    plot_histogram(6, 'I')
  }
})

```



Then pairwise plots for the dynamical parameters (red lines are the true values):

```
withr::with_par(list(mfrow = c(n_examples, 5), mar = c(4, 4, 1, 1)), {
  for (i in seq_len(n_examples)) {
    plot_pair <- function(ind1, ind2, lab1, lab2) {
      plot(
        example_samples$samples[, i, ind1],
        example_samples$samples[, i, ind2],
        xlab = lab1,
        ylab = lab2,
        xlim = range(example_samples$samples[, , ind1]),
        ylim = range(example_samples$samples[, , ind2]),
        main = paste('Example', i)
      )
      abline(v = example_samples$target[i, ind1], col = 'red')
      abline(h = example_samples$target[i, ind2], col = 'red')
    }
    plot_pair(2, 3, 'beta', 'gamma')
    plot_pair(2, 4, 'beta', 'epsilon')
    plot_pair(3, 4, 'gamma', 'epsilon')
    plot_pair(2, 5, 'beta', 'S_t')
    plot_pair(3, 5, 'gamma', 'S')
  }
})
```

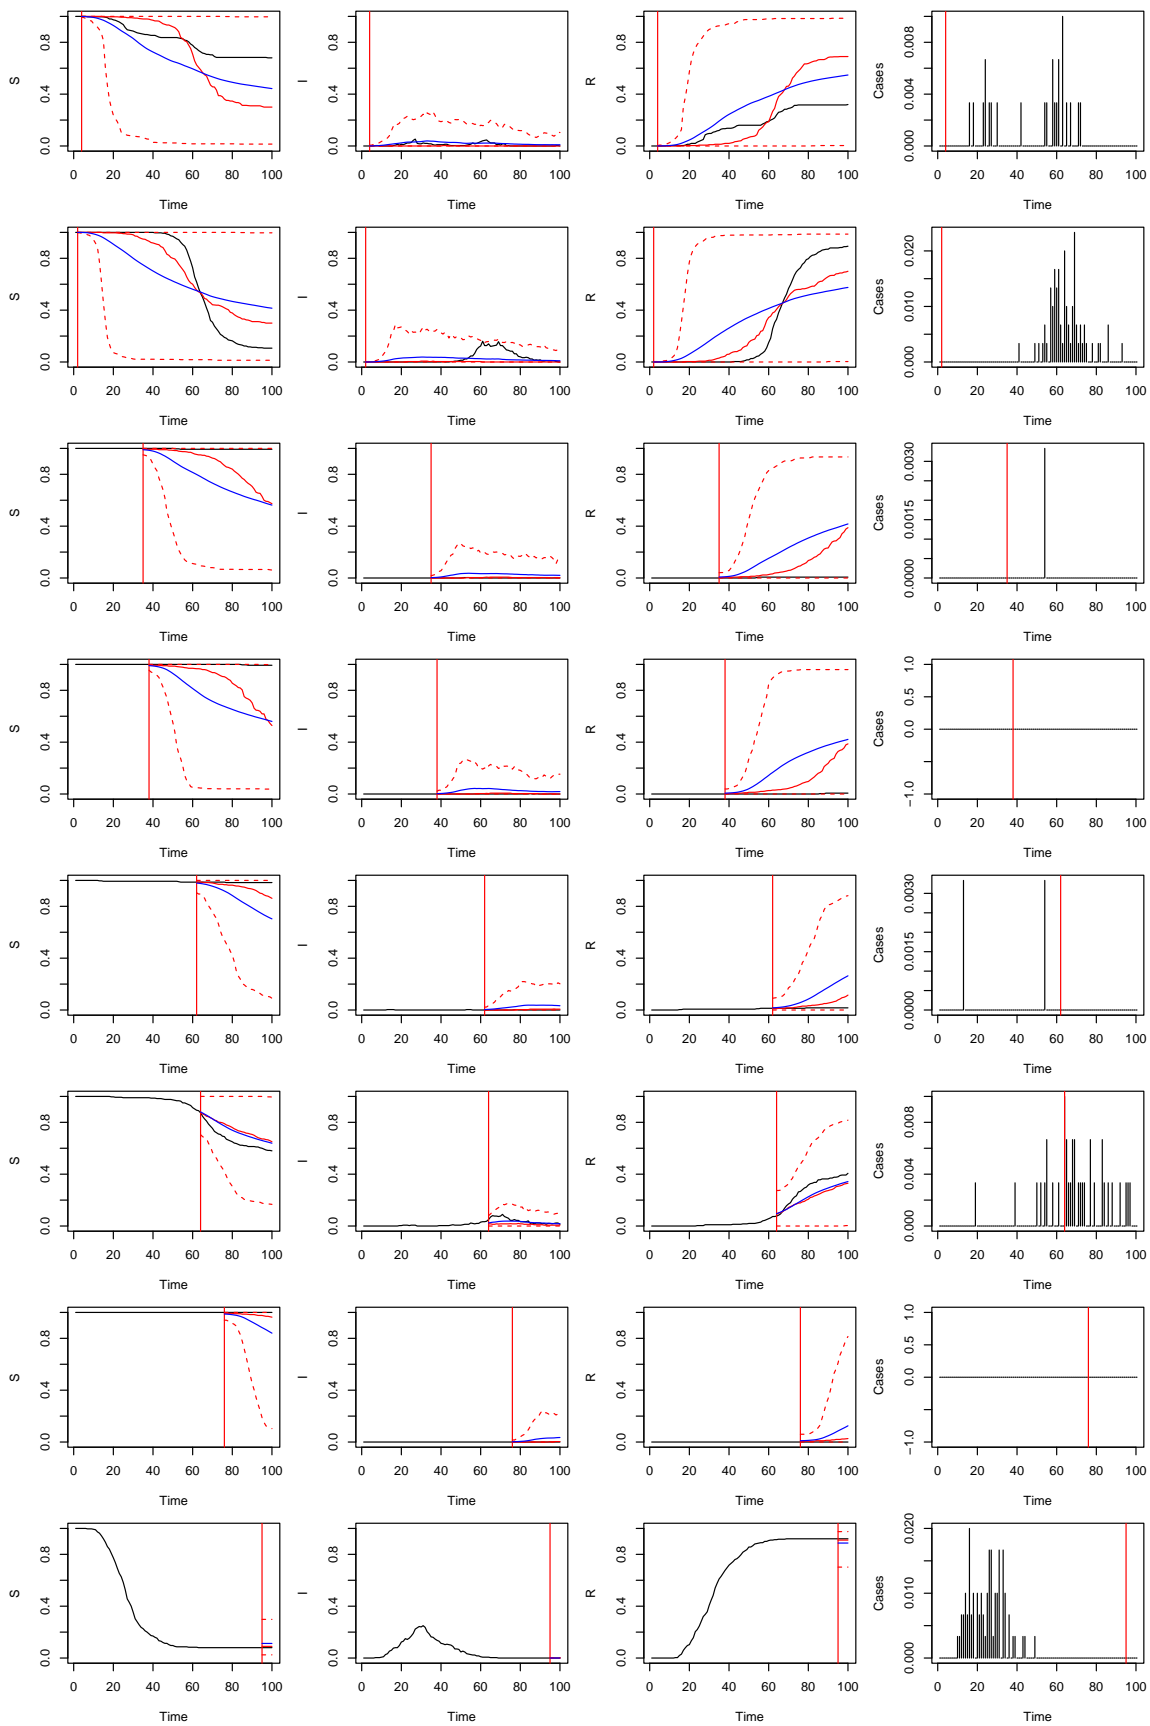


Then we plot the true and forecasted trajectories of the latent variables. In this plot:

- The black line is the true trajectory.
- The red lines are the 2.5th, 50th, and 97.5th posterior percentiles of the forecasted trajectories.
- The blue line is the posterior mean of the forward simulated trajectories.

```
plot_forward_samples <- function(samples) {
  n_examples <- length(samples$forward_samples)
  withr::with_par(list(mfrow = c(n_examples, 4), mar = c(4, 4, 1, 1)), {
    for (i in seq_len(n_examples)) {
      forward_samples_i <- samples$forward_samples[[i]]
      for (j in 1 : 3) {
        plot(
          1 : 100,
          samples$latent[i, , j],
          type = 'l',
          xlab = 'Time',
          ylab = c('S', 'I', 'R')[j],
          ylim = c(0, 1)
        )
        abline(v = samples$t[i], col = 'red')
        forward_mean <- apply(forward_samples_i[[j]], 2, mean)
        forward_quantiles <- apply(forward_samples_i[[j]], 2, quantile, c(0.025, 0.5, 0.975))
        lines(forward_samples_i$time, forward_quantiles[1, ], lty = 2, col = 'red')
        lines(forward_samples_i$time, forward_quantiles[2, ], lty = 1, col = 'red')
        lines(forward_samples_i$time, forward_quantiles[3, ], lty = 2, col = 'red')
        lines(forward_samples_i$time, forward_mean, lty = 1, col = 'blue')
      }
      plot(
        1 : T,
        samples$cases[i, ],
        type = 'h',
        xlab = 'Time',
        ylab = 'Cases',
        ylim = c(0, max(samples$cases[i, ]))
      )
      abline(v = samples$t[i], col = 'red')
    }
  })
}

plot_forward_samples(example_samples)
```



Let's also focus on the single example from the start of this document and step through some forecasts.

```
single_example_times <- seq(10, 90, by = 10)
single_example <- list(
  target = torch_tensor(do.call(rbind, lapply(single_example_times, function(t) {
    c(
      log(0.25),
      log(0.5),
      log(0.2),
      log(0.0001),
      logit(clamp(example_simulation$S[, t], epsilon, 1 - epsilon)),
      logit(clamp(example_simulation$I[, t], epsilon, 1 - epsilon))
    )
  })), device = device),
  conditioning = torch_tensor(do.call(rbind, lapply(single_example_times, function(t) {
    c(t, example_simulation$cases[1, ])
  })), device = device),
  latent = torch_tensor(abind::abind(lapply(single_example_times, function(t) {
    cbind(example_simulation$S[1, ], example_simulation$I[1, ], example_simulation$R[1, ])
  })), along = 0), device = device)
)
single_example_samples <- sample_dataset(256L, single_example)
n_single_examples <- length(single_example_times)

plot_forward_samples(single_example_samples)
```